



CUDA

Parallel programming on pixels

Disclaimer

- ❑ All pictures in this PowerPoint file are taken from the slides by professor Wen-Mei W. Hwu, distributed by NVidia.



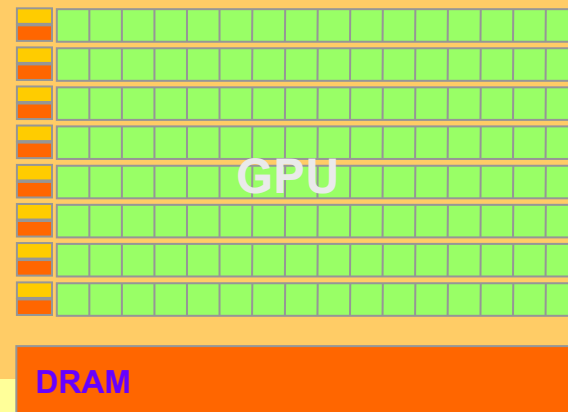
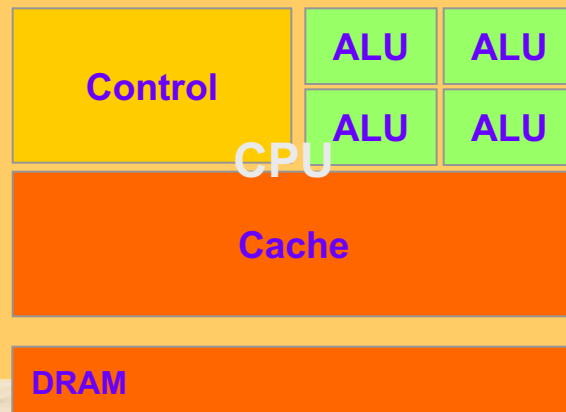
GPU

- ❑ Graphical Processing Unit
- ❑ Used to display 3D graphics on your PC
- ❑ Consists of very a very large number of transistors.
- ❑ Highly parallel computation
- ❑ Cost-effective
 - ❑ The only part that you might consider updating in your PC.

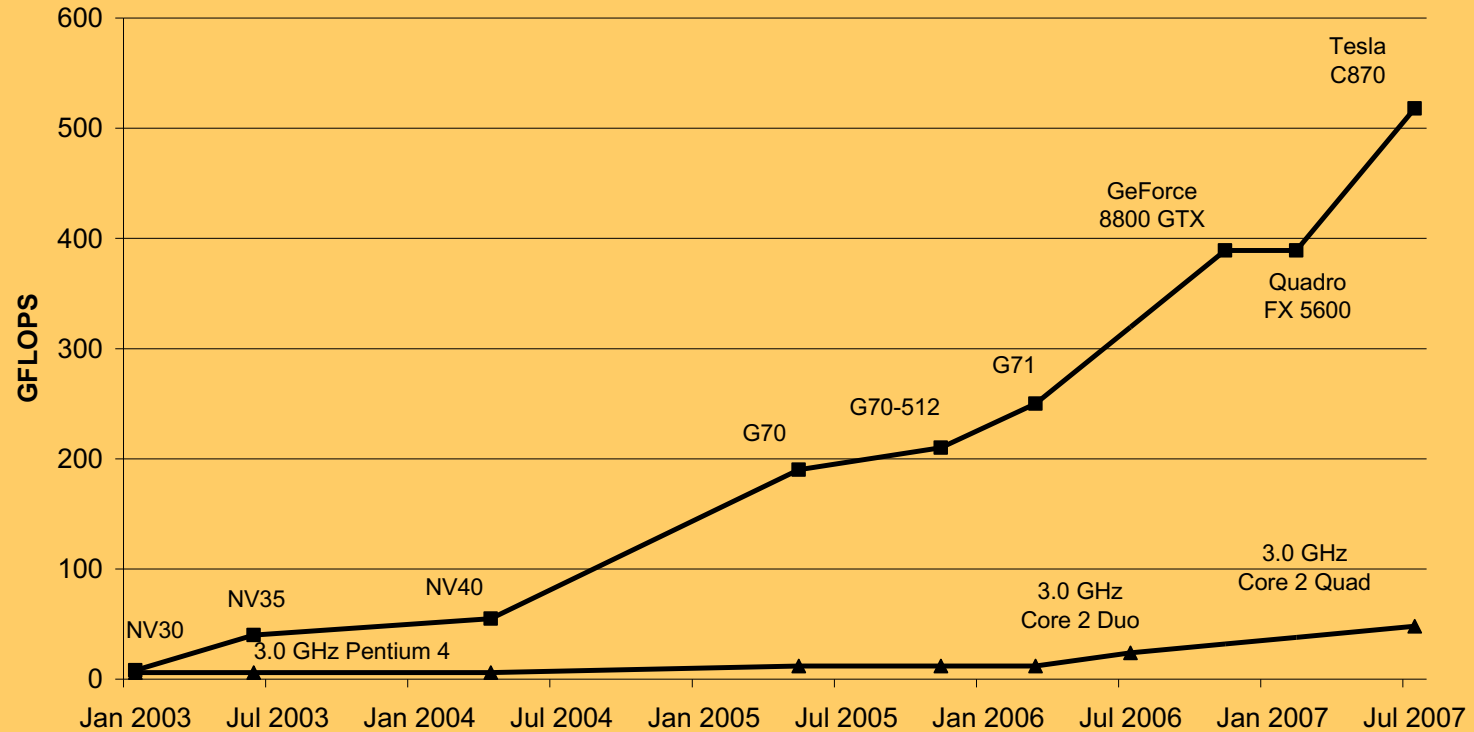


CPU vs. GPU

- ❑ Less data cache and flow control
- ❑ More Computation



Performance in Gflops



GPGPU

□ General Purpose computation using GPU --
not only for pixels.



Advantages

- ❑ Large data arrays, streaming throughput
- ❑ Very large memory bandwidth
- ❑ Fine-grain SIMD parallelism
 - ❑ A very large number of threads
- ❑ Low-latency floating point (FP) computation
 - ❑ High power computation capability
- ❑ Piggyback on the fast advancing GPU technology



GPGPU Applications

- ❑ Tons of application on GPGPU.org and Nvidia websites.
- ❑ http://www.nvidia.com/object/cuda_home.html#
 - ❑ Game effects (FX) physics
 - ❑ image processing
 - ❑ Physical modeling
 - ❑ computational engineering
 - ❑ matrix algebra
 - ❑ Convolution
 - ❑ correlation, sorting



Implication

- ❑ Suitable for data-parallel processing
 - ❑ The same computation is performed on many data elements in parallel
 - ❑ Low control flow overhead
 - ❑ high floating point arithmetic intensity
- ❑ Computation intensive threads on large number of data hide large memory latency for each
 - ❑ No need for large data cache



Why not before?

- ❑ Graphics API only
- ❑ Addressing modes is limited by texture size/dimension
- ❑ Limited outputs due to shader capabilities
- ❑ No instruction on integer and bits
- ❑ No interaction between pixels



CUDA in a Long Sentence

- ❑ Compute Unified Device Architecture
- ❑ A general purpose parallel computing architecture that leverages the parallel compute engine in NVIDIA graphics processing units (GPUs) to solve many complex computational problems in a fraction of the time required on a CPU.



CUDA in a Short Sentence

- ❑ Integrates CPU+GPU application C program by running serial parts on CPU and highly parallel parts on GPU.



CUDA in Your PC

- ❑ Can run on your PC with NVIDIA display card.
 - ❑ NVIDIA CUDA-enabled products
 - ❑ http://www.nvidia.com/object/cuda_learn_products.html
- ❑ CUDA download
 - ❑ http://www.nvidia.com/object/cuda_get.html



An Illustration

CPU Serial Code

GPU Parallel Kernel

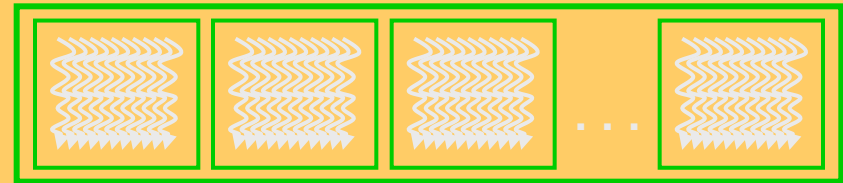
```
KernelA<<< nBlk, nTid >>>(args);
```

CPU Serial Code

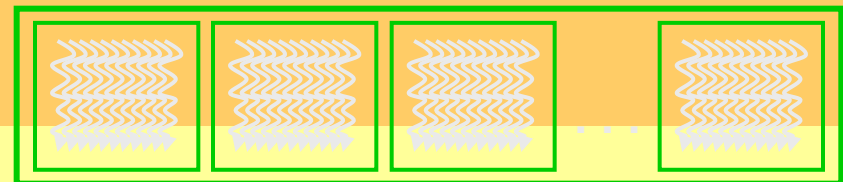
GPU Parallel Kernel

```
KernelB<<< nBlk, nTid >>>(args);
```

Grid 0



Grid 1



Terminology

- ❑ Host
 - ❑ CPU
- ❑ Device
 - ❑ GPU
- ❑ Kernel
 - ❑ Computation intensive C functions running on devices as threads.
- ❑ Grid
 - ❑ The view of cores within a device
- ❑ Block
 - ❑ A basic unit in a device



The Programming Model

- ❑ The *host* runs the *sequential* part of the application.
- ❑ The *parallel* data intensive parts are written as *kernel* functions, and sent to *devices* for execution (as *threads*).
- ❑ The *blocks* are the unit for computation and they are organized as a *grid*.



C Program Sequential Execution

Serial code

Parallel kernel
Kernel0<<<>>>()

Serial code

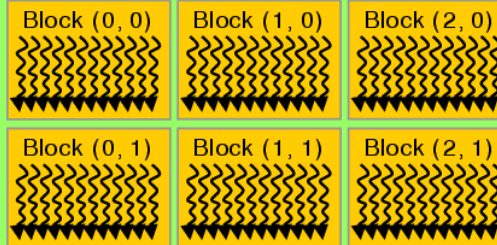
Parallel kernel
Kernel1<<<>>>()

Host



Device

Grid 0

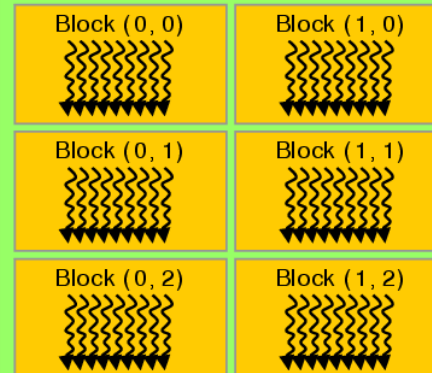


Host



Device

Grid 1



GPU and CPU Thread

- ❑ GPU threads are very lightweight and have very little creation overhead, but CPU thread creation overhead is very high.
- ❑ The implication is that GPU can use a huge number of threads to achieve high performance, but CPU can only use a relatively small number of threads.



Thread id

- ❑ Each thread has an id so that it will be able to determine which part of data it is supposed to handle.
- ❑ The thread id is stored in a variable threadIdx
- ❑ The contents of threadIdx depends on the dimensionality of the grid.



Thread Blocks

- ❑ A block is an N-dimensional array of threads.
 - ❑ N could be 1, 2, or 3.
- ❑ The coordinates of a thread within a block are stored within the variable `threadId` as a vector of three elements, as x, y, and z.
- ❑ When the host calls the kernel function to run on the devices, it can specify the thread block layout between the name of the kernel function and the parameter list.
 - ❑ `initialize <<<1, 5>>> (int_array);`



The ThreadIdx

- ❑ The kernel function can retrieve the thread index by referencing threadIdx variable.
- ❑ `Int_array[threadIdx.x] = threadIdx.x;`
- ❑ Note that this is for a one-dimensional layout, so we use the x element only.



Global Memory

- ❑ The host and devices can share memory.
- ❑ This special global memory must be allocated using special allocation routine. Note that the prototypes of `cudaMalloc` is different from `malloc`.
 - ❑ `cudaMalloc(&pointer, size);`
 - ❑ The first argument is the address of the pointer.
 - ❑ `cudaFree(pointer);`



Host Memory

- ❑ The host can allocate and free its own memory using malloc and free.



CUDA Computation

- ❑ The host allocates and initializes its memory.
- ❑ The host allocates the global memory.
- ❑ The host copies the host memory into global memory.
- ❑ The host calls kernel function to perform computation on the global memory.
- ❑ The host copies the global memory back to the host memory.
- ❑ The host outputs the results.



Memory Transfer

- ❑ Host to device

- ❑ `cudaMemcpy(device_memory, host_memory, cudaMemcpyHostToDevice);`

- ❑ Device to host

- ❑ `cudaMemcpy(host_memory, device_memory, size, cudaMemcpyDeviceToHost);`

- ❑ Note that it is always from the second argument to the first argument.

