



CUDA

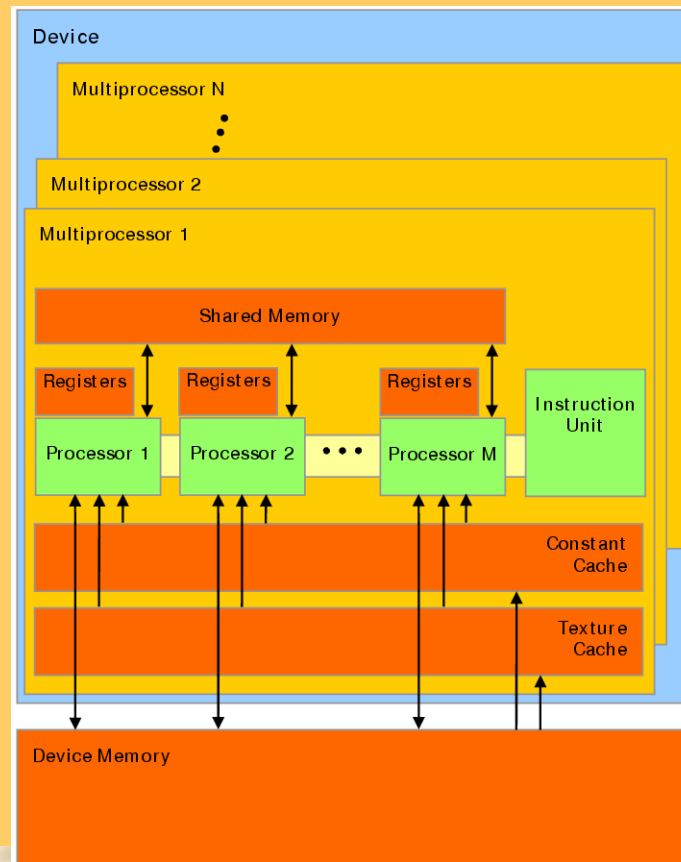
Parallel programming on pixels

Hardware

- ❑ A GPU has many multiprocessors.
- ❑ A multiprocessor has
 - ❑ Core for computation
 - ❑ Register for computation
 - ❑ Memory for storage



An Illustration of GPU



A set of SIMT multiprocessors with on-chip shared memory.

Figure 3-1. Hardware Model

Execution

- ❑ Kernel function is run in grid.
 - ❑ `<<<grid-dim, block-dim>>>`
- ❑ A grid has many blocks.
 - ❑ The layout of the grid is specified by grid-dim.
 - ❑ The layout of the block is specified by block-dim.



Block and Multiprocessor

- ❑ A thread block runs on a multiprocessor as a unit.
 - ❑ No “partial” blocks
 - ❑ No migration
- ❑ A multiprocessor may run multiple blocks.
 - ❑ A multiprocessor has the resource to run many threads simultaneously.
 - ❑ The number of threads is limited by resource. If one wish to have more threads, he should get more blocks, not more threads per block.



Synchronization

- ❑ Threads with the same block can synchronize with `__syncthreads()`.
- ❑ Kernel waits for all previous CUDA calls returns, but the control returns to host immediately.
- ❑ `cudaMemcpy()` starts after all previous CUDA calls returns, and return only when memory operation finishes.



Machine Configuration

- ❑ Use `cudaGetDeviceCount(&device_count)` to determine the number of devices in the system.
- ❑ Use `cudaGetDeviceProperties(cudaDeviceProp *, int deviceid)` to determine the important parameters in performance tuning.
 - ❑ Warp size
 - ❑ Maximum threads per block



cudaDeviceProp

- ❑ `char name[256];`
- ❑ `size_t totalGlobalMem;`
- ❑ `size_t sharedMemPerBlock;`
- ❑ `int regsPerBlock;`
- ❑ `int warpSize;`
- ❑ `size_t memPitch;`
- ❑ `int maxThreadsPerBlock;`
- ❑ `int maxThreadsDim[3];`
- ❑ `int maxGridSize[3];`



cudaDeviceProp

- ❑ `size_t totalConstMem;`
- ❑ `int major;`
- ❑ `int minor;`
- ❑ `int clockRate;`
- ❑ `size_t textureAlignment;`
- ❑ `int deviceOverlap;`
- ❑ `int multiProcessorCount;`



Function Qualifiers

- ❑ `__global__`

- ❑ Host code invokes, runs on device.

- ❑ `__device__`

- ❑ Device code invokes, runs on device.

- ❑ `__host__`

- ❑ Host code invokes, runs on host.



__global__ Restrictions

- ☐ Must return void.
- ☐ No recursion
- ☐ No static variables
- ☐ No variable number of arguments



Combination

- ❑ `__host__` and `__device__` can be combined.
 - ❑ The same code could be run on both device and host for performance and convenience.
 - ❑ Same source but different binary codes for CPU and GPU.



Built-in Device Variables

- ❑ All of type dims
 - ❑ With constructor, e.g. (N, N).
 - ❑ Has component x, y, and z.
- ❑ Block id and grid dimension
 - ❑ gridDim
 - ❑ blockIdx
- ❑ Thread id and block dimension
 - ❑ threadIdx
 - ❑ blockDim

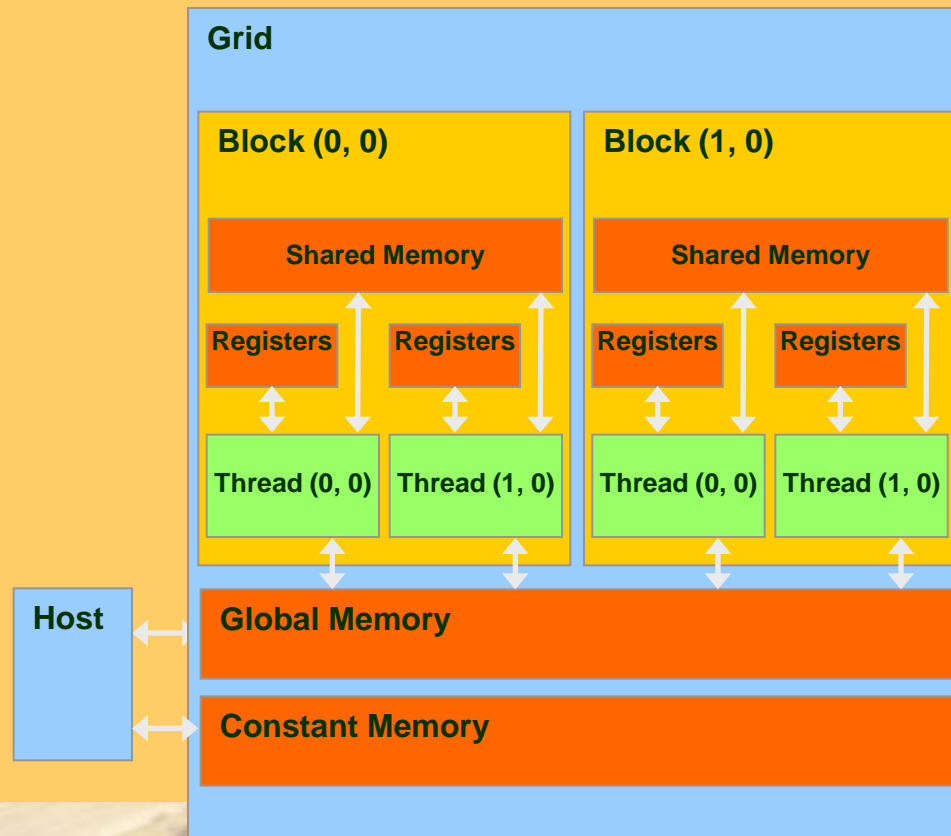


Memory Hierarchy

- ❑ Global and constant memory
 - ❑ Large
 - ❑ Shared by all threads and host
 - ❑ Off-chip
- ❑ Shared memory
 - ❑ Small
 - ❑ Shared by all threads in the same block
 - ❑ On-chip



An Illustration



© David Kirk/NVIDIA and Wen-mei W. Hwu
Taiwan, June 30-July 2, 2008

Implications

- ❑ A large number of threads hide instruction and memory latency for each other.
- ❑ Threads of the same block can share data via fast shared memory.
- ❑ All threads can share data via global memory.



GPU Variable Types

☐ Classification

- ☐ `__device__`

- ☐ `__shared__`

- ☐ `__constant__`

- ☐ `__local__`



GPU Variable Types

- ❑ Where the variables will be allocated
- ❑ The lifetime of the variable



__device__ Variables

- ❑ So-called global memory
- ❑ Large and slow
- ❑ Allocated with cudaMalloc
- ❑ Accessible by all threads and host
- ❑ Valid throughout the entire execution



__shared__ Variables

- ❑ So called shared memory
- ❑ Small and fast
- ❑ Allocated by declaration or kernel invocation
- ❑ Accessible by threads in the same block, but values assigned are guaranteed to be visible by other threads only after `__syncthreads()`.
- ❑ Valid only during kernel execution



__shared__ by Declaration

```
__global__ void kernel(...)  
{  
    __shared__ float sData[256];  
    ...  
}  
  
int main(void)  
{  
    kernel<<<griddim, blockdim>>>(...);  
}
```



__shared__ by Kernel Invocation

```
__global__ void kernel(...)  
{  
    extern shared float sData[]; /* we will have 256 floats here. */  
    ...  
}  
  
int main(void)  
{  
    kernel<<<griddim, blockdim, 256 * sizeof(float)>>>(...);  
}
```



Kernel Invocation

- ❑ <<<griddim, blockdim, sharedmemsize>>>
- ❑ sharedmemsize is optional, and the default value is 0.
- ❑ If not 0 then it specifies the number of bytes of shared memory allocated for each block.
 - ❑ The allocate shared memory can be accessed as sh_data.



__constant__ Variables

- ❑ Stored in constant memory
- ❑ Read-only on device, but can be set on host.
- ❑ Accessible by all threads and host
- ❑ Valid throughout the entire execution



__local__ Variables

- ❑ Stored in registers if not array
- ❑ Allocated by declaration
- ❑ Accessible by the thread that declares it.
- ❑ Valid when the thread is active



Access

- ☐ Host may access

- ☐ `__global__`

- ☐ `__constant__`

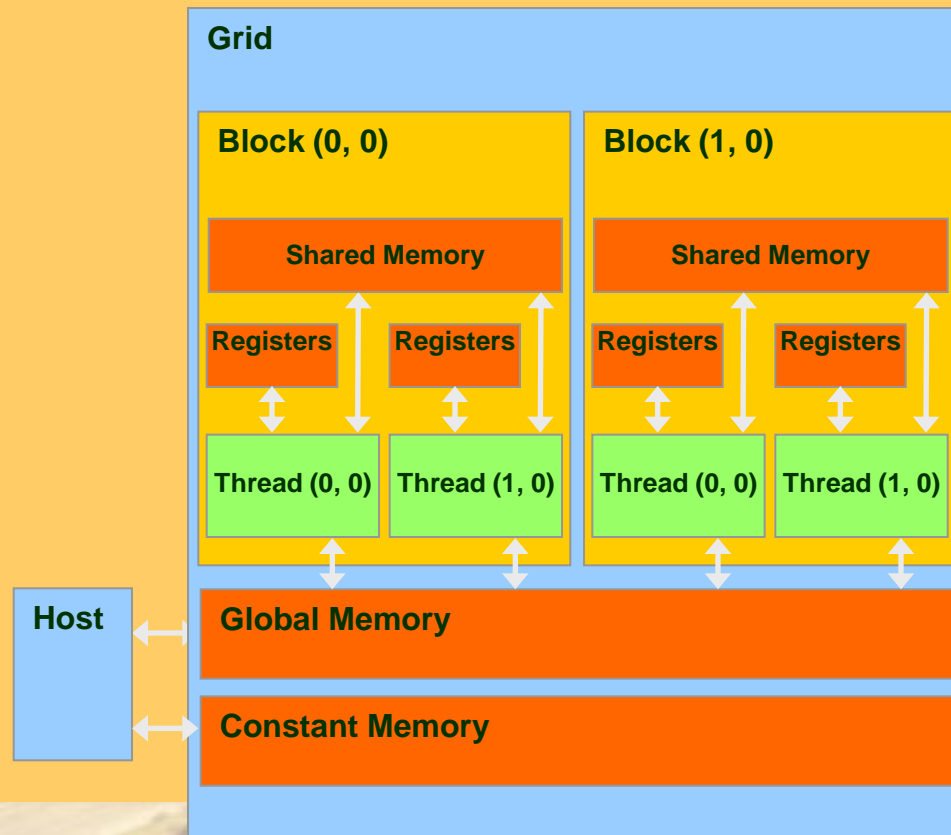
- ☐ Kernel only

- ☐ `__local__`

- ☐ `__shared__`



An Illustration



© David Kirk/NVIDIA and Wen-mei W. Hwu
Taiwan, June 30-July 2, 2008

Restriction

- ❑ No qualifier on
 - ❑ Members of a structure
 - ❑ Host local variables
- ❑ `__shared__` cannot be initialized during declaration.



Strategy

- ❑ Shared memory is much faster than global memory
 - ❑ Host moves data into global memory.
 - ❑ Thread moves data from global memory into private memory for processing.
 - ❑ Thread computes.
 - ❑ Copy results from private to global memory.



Timing

- ❑ The `clock()` routine returns the wall clock time in clock ticks as a `clock_t`.
- ❑ The number of clock ticks per second is in the constant `CLOCKS_PER_SEC`.
- ❑ Before calling `clock` make sure all CUDA routines finish by calling `cudaThreadSynchronize()` at host.



Thread Synchronization

- ❑ `__syncthreads()`
 - ❑ Calls from device code
 - ❑ Synchronizes all threads in the same block
 - ❑ Enforces consistent view on shared memory
- ❑ `cudaThreadSynchronize()`
 - ❑ Calls from host code.
 - ❑ Synchronizes all threads
 - ❑ Enforces correct timing on kernel



Matrix Multiplication

- ❑ Two versions to compute $C = A \times B$
- ❑ Both A and B are in global memory
 - ❑ Version 1 retrieves data from global memory directly.
 - ❑ Version 2 retrieves data from global memory and stores in shared memory.



Using Shared Memory

- ❑ Organize the grid as a N/b by N/b grids, and each block has b by b threads.
- ❑ Each thread block will declare two b by b shared arrays in shared memory (for A and B).
- ❑ Each thread in a block will be responsible for loading an element from A , and an element from B .





Matrix Loading

- ❑ The matrix A and B will be divided into $N/b * N/b$ sub-matrices, each of size b by b .
- ❑ If a block is in the i -th row and j -th column of the grid, then the threads in the block will load the first sub-matrix (of size b by b) from the i -th sub-matrix row of A , and the first sub-matrix (of size b by b) from the j -th sub-matrix column of B , then the threads will load from the second sub-matrix, and so on.



An Illustration

□ First step




A

B

An Illustration


- ❑ Each thread moves an element in A.



A in global memory

An Illustration


- Each thread moves an element in B




B in global memory

An Illustration

□ Second step





A



B

An Illustration

□ Third step





A

B

An Illustration

□ Final step



A

B


Inner Product

- ❑ Every time the threads load two sub-matrix (A and B), it performs an inner product it is responsible for.
 - ❑ The thread at the i -th-row and j -th column of a thread block will compute the inner product of the i -th row of sub-matrix A and j -th column of sub-matrix B.
- ❑ The sum of these inner products will be the final C_{ij}



An Illustration

□ Each thread computes an element in C



Comparison

- ❑ Each thread of version 1
 - ❑ Load $2N$ element (A, B) from global memory for inner product.
- ❑ Each thread of version 2
 - ❑ Load $2(N / b)$ element (A, B) from global memory.
 - ❑ Load $2N$ elements from shared memory (for inner product).
 - ❑ The threads “share” the loading from slow global memory.

