



CUDA

Parallel programming on pixels

Performance Tuning

- ☐ Machine Configuration
- ☐ Accurate Timing
- ☐ Block size
- ☐ Control flow
- ☐ Memory access
- ☐ Unroll
- ☐ Double buffering



Machine Configuration

- ❑ Use `cudaGetDeviceCount(&device_count)` to determine the number of devices in the system.
- ❑ Use `cudaGetDeviceProperties(cudaDeviceProp *, int deviceid)` to determine the important parameters in performance tuning.
 - ❑ Warp size
 - ❑ Maximum threads per block



cudaDeviceProp

- ❑ `char name[256];`
- ❑ `size_t totalGlobalMem;`
- ❑ `size_t sharedMemPerBlock;`
- ❑ `int regsPerBlock;`
- ❑ `int warpSize;`
- ❑ `size_t memPitch;`
- ❑ `int maxThreadsPerBlock;`
- ❑ `int maxThreadsDim[3];`
- ❑ `int maxGridSize[3];`



cudaDeviceProp

- ❑ `size_t totalConstMem;`
- ❑ `int major;`
- ❑ `int minor;`
- ❑ `int clockRate;`
- ❑ `size_t textureAlignment;`
- ❑ `int deviceOverlap;`
- ❑ `int multiProcessorCount;`



Acurate Timing

- ❑ Use cuda events for accurate timing, since event has time information.
- ❑ Place start event into the event stream 0.
- ❑ Place the code for timing.
- ❑ Place end event into the event stream 0.
- ❑ Wait for the end event to finish.
- ❑ Compute the elapse time between start and end.



Event Routines

- ❑ Declared as type `cudaEvent_t`
- ❑ Created using `cudaEventCrreate(*cudaEvent_t)`
- ❑ Record into the operation stream by `cudaEventRecord(cudaEvent_t, int stream)`.
 - ❑ Use stream 0, which is for all events.
- ❑ Synchronized with event using `cudaEventSynchronize(cudaEvent_t)`.
- ❑ Compute the elapsed time using `cudaEventElapsedTime(float *elapsed_time, cudaEvent_t start, cudaEvent_t end)`.
- ❑ Destroyed with `cudaEnentDestroy(cudaEvent_t)`.



An Example

```
float time;  
cudaEvent_t start, stop;  
cudaEventCreate(&start);  
cudaEventCreate(&stop);
```



An Example

```
cudaEventRecord(start, 0);  
multiply <<< grids, blocks >>> ((int (*)(N))device_A, (int  
    (*)(N))device_B, (int (*)(N))device_C);  
cudaEventRecord(stop, 0);  
cudaEventSynchronize(stop);  
cudaEventElapsedTime(&time, start, stop);  
printf("the multiplication takes %f seconds\n", time);  
cudaEventDestroy(start);  
cudaEventDestroy(stop);
```



Block Size

- ❑ How many threads should go into a block?
- ❑ Limitation on the resource of multiprocessors



Thread Block

- ❑ Up to 512 threads per block
- ❑ Up to three dimensions
- ❑ Has shared memory
- ❑ Can synchronize with `__syncthread()`



Multiprocessor

- ❑ So called Streaming Multiprocessors (SM)
- ❑ A SM can run a limited number of blocks and threads.
 - ❑ 8 blocks and 768 threads for G80
- ❑ SM will maintain thread and block id for all thread blocks running on it, and schedule these threads for execution.



Warps

- ❑ Thread are scheduled in unit of 32 (called Warp).
- ❑ CUDA Programmer cannot see this, but should be aware of this.
- ❑ Only one warp (from all blocks assigned to a SM) will execute on a SM at any given time.



Warp Scheduling

- ❑ Operand-ready warps are eligible for execution.
- ❑ It takes 4 cycles to dispatch an instruction to a warp.
- ❑ It takes 200 cycles to fetch from global memory.
- ❑ How many warps does it take to hide memory latency?




Block Size

- ❑ 8 by 8 block size
 - ❑ $768 / 64 = 12 > 8$ blocks
 - ❑ Only 8 blocks = 512 threads are used
- ❑ 16 by 16 blocks
 - ❑ $768 / 256 = 3$ blocks
 - ❑ All 3 blocks = 768 threads are used
- ❑ 32 by 32
 - ❑ 1024 threads cannot go into a SM



Tiling Size Effects



© David Kirk/NVIDIA and Wen-mei W. Hwu, Taiwan, June 30-July 2, 2008

Control Flow

- ❑ SPMD program

- ❑ Threads do things according to thread id.

- ❑ A warp should follow the same control path as often as possible.



Granularity of Control Flow

- ❑ If threads with a warp take different control paths, they will be serialized.
- ❑ Remember that the same instruction will be issued to all threads in the same block.
- ❑ That means we will not have good speedup.
- ❑ Therefore we should get all threads in a warp to go along the same control flow as much as possible.



Examples

- ❑ `if (threadIdx.x > 2)`

- ❑ Thread 0, 1, and 2 will take else.

- ❑ 30 other threads take then.

- ❑ Two groups will be serialized.

- ❑ `If (threadIdx.x / groupsize > 2)`

- ❑ If groupsize is a multiple of 32, all threads in a warp go along the same control path.




Sum of N Numbers

- ❑ To add the numbers in a global array of N elements.
- ❑ Use a block of N threads to compute the sum.
- ❑ The computation is in $\log N$ phases.
 - ❑ In the first phase all even number threads with index i add the element in the odd number threads that have indices $i + 1$ to them.
 - ❑ The next phase all threads of indices multiple of 4 add the elements in the threads that have indices $i + 2$ to them.
 - ❑ We repeat this process until the first element has the sum of all elements.



An Illustration



© David Kirk/NVIDIA and Wen-mei W. Hwu, Taiwan, June 30-July 2, 2008


Reverse Order

- ❑ Add the element in decreasing strides.
- ❑ In the first iteration, those threads with indices smaller than $N/2$ will add the elements that have indices $N/2$ greater than them to the element they have.
- ❑ Then the stride reduces by half and we do it again.
- ❑ We repeat the process until the first element has the sum of all numbers.



An Illustration

Thread 0



© David Kirk/NVIDIA and Wen-mei W. Hwu, Taiwan, June 30-July 2, 2008

Memory Access

- ❑ A warp should access memory in horizontal (contiguous) order .
- ❑ If horizontal access is not possible, one should reduce the stride in rows.




Tiled Matrix Multiplication

- ❑ In version 1 (using global memory only), the warps reference data horizontally (N) and vertically (N).
- ❑ In version 2 (using global and shared memory), the warps reference data horizontally (b) and vertically (b).




Tiling Size Effects




© David Kirk/NVIDIA and Wen-mei W. Hwu, Taiwan, June 30-July 2, 2008

An Illustration

□ First step




A



B

An Illustration


- ❑ Each thread moves an element in A.



A in global memory

An Illustration

- Each thread moves an element in B




sB in shared memory

B in global memory


An Illustration

- Each thread computes an element in C using shared memory.



An Illustration

- ❑ Each thread computes an element in C using global memory.



Unroll

- ❑ Inline the most time consuming loop.
- ❑ `for (k = 0; k < b; k++)`
 - ❑ `sum += sA[threadIdx.x][k] * sB[k][threadIdx.y];`
- ❑ Change to the following.
 - ❑ `sum += sA[threadIdx.x][0] * sB[0][threadIdx.y]`
`+ sA[threadIdx.x][1] * sB[1][threadIdx.y] + ... +`
`sA[threadIdx.x][7] * sB[7][threadIdx.y];`



Advantage

- ❑ Without loop overhead
 - ❑ All indices are determined at compile time.
- ❑ Hopefully better function unit utilization.
- ❑ Larger amount of computation between jumps.



Double Buffering

- ❑ Use two buffers
 - ❑ One for ongoing memory fetch.
 - ❑ One for current computation.



An Illustration

- ❑ Load A and B into “current” shared memory
- ❑ Repeat
 - ❑ Load A and B into “next” shared memory.
 - ❑ Compute on “current” shared memory.
 - ❑ Switch the current and next memory.

