Advanced OpenCL Programming

Pangfeng Liu National Taiwan University

June 7, 2015



Multiple Devices

- We have been using one device for computation.
- Now we will use multiple device to solve a problem.
- We will use matrix multiplication as an example.

Partition

- Since we are very poor and cannot afford four GPUs, we will partition the data by rows (not by block) among devices.
- We will partition matrix by rows, so that the memory assigned to a device is contiguous.
- Each kernel will run on a device, and compute part of the answers.

Partition

- We will partition A and C by rows, and do *not* partition B.
- For example, if we use two devices, then the first kernel will compute the top half of *C*, and another kernel will compute the bottom half of *C*.
- The reason for not partitioning *B* is that all kernels need the entire *B*.

Constants

- MAXLOG is the maximum number of bytes for storing compilation log. More on this later.
- The number of device is DEVICENUM. We will use 2 in our humble installation.
- ITEMPERDEVICE is the number of work item in NDRange of a kernel.

Constants

Example 1: (matrixMul-time-copy-local-multidevice.c)

```
#define CL_USE_DEPRECATED_OPENCL_2_O_APIS
   #include <stdio.h>
   #include <assert.h>
4
   #include <CL/cl.h>
5
6
   #define N 1024
   #define Blk 64
   #define BSIDE (N / Blk)
9
   #define MAXGPU 10
10
   #define MAXK 1024
11
   #define MAXI.OG 4096
12
   #define DEVICENUM 2
13
   #define ITEMPERDEVICE
                           (N * N / DEVICENUM)
14
   #define NAND2SECOND 1000000000.0
15
16
   cl_uint A[N][N], B[N][N], C[N][N];
```

Discussion

• If *N* is 1024, and the number of device is 2, then how many works items are there in a kernel?

Select Devices

- We will select the first DEVICENUM GPU device from our humble installation.
- The code is similar to those earlier version, except that now we need to make sure that the number of GPUs found is at least DEVICENUM.

Devices

Example 2: (matrixMul-time-copy-local-multidevice.c)

```
29
     cl_device_id GPU[MAXGPU];
30
     cl_uint GPU_id_got;
31
     status =
32
        clGetDeviceIDs(platform_id, CL_DEVICE_TYPE_GPU,
33
                       MAXGPU, GPU, &GPU_id_got);
34
     assert(status == CL_SUCCESS &&
35
             GPU_id_got >= DEVICENUM);
36
     printf("There are %d GPU devices\n", GPU_id_got);
```

Context

- Since a context can consist of multiple devices, we need only one context.
- Here we include the first DEVICENUM GPUs in the context.

Context

Example 3: (matrixMul-time-copy-local-multidevice.c)

Multiple Queue

- Recall that a command queue connects to a device.
- Since we have multiple devices, we need multiple command queues.
- We put the command queues in the commandQueue array.
- Note that we set the CL_QUEUE_PROFILING_ENABLE to enable profiling.

Command Queue

Example 4: (matrixMul-time-copy-local-multidevice.c)

One Context, Multiple Devices

- Up to this point, we can image there is only one context, which has multiple devices.
- Within this context we will have one command queue to connect to each device.
- Later we will send commands to these devices. The command sent into a command queue will run on the device this command queue connects to.

Discussion

• What will be sent into these command queues as commands?

Program

- The kernel will be sent into the command queue as commands.
- Before this can happen we need to compile the kernel.
- Recall that the "kernel" at this point is only a set of strings, which need to be compiled.

Build Program

- We call clBuildProgram to build executable.
- Note that we pass all devices we want to use as parameters, so
 OpenCL will build executable for all our DEVICENUM devices.

Compilation Log

- The kernel source is compiled when an OpenCL program runs.
- It is very inconvenient because when we run an OpenCL program the execution will not display compilation errors of the kernel.
- We will use clGetProgramBuildInfo to show the error message if the compilation of the kernel fails.

Prototype 5: clGetProgramBuildInfo.h

Parameters

```
program The compiled program.
  device The device you want to query.
param_name The information you want to query.
param_value_size The length of param_value in bytes.
param_value The location for the query answer.
param_value_size_ret The number of bytes returned from the query.
```

Error

- When there is an error in clBuildProgram, we will call clGetProgramBuildInfo to find out.
- Since we are not sure about which device causes the compilation error, we list the information from all of them.
- We query with CL_PROGRAM_BUILD_LOG for the compilation log, and place it into the log buffer we prepared.

Build Program

Example 6: (matrixMul-time-copy-local-multidevice.c)

```
64
      status =
65
        clBuildProgram (program, DEVICENUM, GPU, NULL,
66
                        NULL. NULL):
67
      if (status != CL_SUCCESS) {
68
        char log[MAXLOG];
69
        size_t logLength;
70
        for (int device = 0; device < DEVICENUM; device++)</pre>
71
          clGetProgramBuildInfo(program, GPU[device],
72
                                  CL_PROGRAM_BUILD_LOG,
73
                                  MAXLOG, log, &logLength);
74
          puts(log);
75
76
        exit(-1);
77
78
      printf("Build program completes\n");
```

Discussion

- What is the type of the compilation log?
- What does puts do?

Partition

- Now we are ready to partition host buffers A, B and C.
- We will create DEVICENUM buffers for A, and each of which will be given to a kernel as a parameter for computations.
- We partition A into sub-matrix of N / DEVICENUM rows of A each, and assign each partition to an OpenCL buffer.

Buffer A

- ITEMPERDEVICE means the number of items per device.
- The index device determines the starting position of a buffer, i.e., where the matrix A will be partitioned.

Buffer for A

90

91

92

93

94

95

96

97

98

99

Example 7: (matrixMul-time-copy-local-multidevice.c)

```
cl_mem bufferA[DEVICENUM];
for (int device = 0; device < DEVICENUM; device++) {
  bufferA[device] =
    clCreateBuffer(context,
        CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR,
        ITEMPERDEVICE * sizeof(cl_uint),
        ((cl_uint *)A) + device * ITEMPERDEVICE,
        &status);
  assert(status == CL_SUCCESS);
}</pre>
```

Buffer B

• We do not need to partition *B* because it will be used by all kernels.

Buffer for B

101

102

103

104

105

Example 8: (matrixMul-time-copy-local-multidevice.c)

```
cl_mem bufferB =
    clCreateBuffer(context,
        CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR,
        N * N * sizeof(cl_uint), B, &status);
assert(status == CL_SUCCESS);
```

Buffer C

- Buffers for C are similarly built as for A.
- Unlike A and B that will copy host buffers to device buffers,
 C will use host memory directly.

Buffer for C

107

111

Example 9: (matrixMul-time-copy-local-multidevice.c)

```
cl_mem bufferC[DEVICENUM];
108
      for (int device = 0; device < DEVICENUM; device++) {</pre>
109
         bufferC[device] =
110
           clCreateBuffer(context,
             CL_MEM_WRITE_ONLY | CL_MEM_USE_HOST_PTR,
112
             ITEMPERDEVICE * sizeof(cl uint).
113
             ((cl_uint *) C) + device * ITEMPERDEVICE,
114
             &status);
115
         assert(status == CL_SUCCESS);
      }
116
      printf("Build buffers completes\n");
117
```

Discussion

 Make sure that you understand the offset calculation in A and C.

NDRange

- According to our partition we declare NDRange as a two dimension array with N / DEVICENUM rows and N columns.
- This is consistent with C, where each work item will compute an element of C.
- The work group is still BSIDE by BSIDE since we will use the previous local memory algorithm.
- We will have DEVICENUM kernels, so we will need DEVICENUM events for synchronization and profiling.

NDRange

Example 10: (matrixMul-time-copy-local-multidevice.c)

```
size_t globalThreads[] =

{(size_t)(N / DEVICENUM), (size_t)N};

size_t localThreads[] = {BSIDE, BSIDE};

cl_event events[DEVICENUM];
```

Set Argument

- We will launch DEVICENUM kernels. Each of them will compute N / DEVICENUM rows and N columns of C.
- We loop through all devices, and assign the argument order as A, B, and C.
- Note that we need to supply the correct A and C since there
 are DEVICENUM of them.

Set Arguments

Example 11: (matrixMul-time-copy-local-multidevice.c)

```
124
      for (int device = 0; device < DEVICENUM; device++) {</pre>
125
         status = clSetKernelArg(kernel, 0, sizeof(cl_mem),
126
                                  (void*)(&bufferA[device]));
127
         assert(status == CL SUCCESS):
128
         status = clSetKernelArg(kernel, 1, sizeof(cl_mem),
129
                                  (void*)&bufferB):
130
        assert(status == CL SUCCESS):
131
         status = clSetKernelArg(kernel, 2, sizeof(cl_mem),
132
                                  (void*)(&bufferC[device])):
133
         assert(status == CL_SUCCESS);
        printf("Set kernel arguments completes\n");
134
```

Kernel

- Now we are ready to launch kernel.
- Note that we need to supply the command queue to the device as a parameters.
- The same kernel is launched again and again for DEVICENUM times. The only difference is the buffers A and C used in launching the kernel.
- We associate an event with each kernel launching.

Start Kernel

Example 12: (matrixMul-time-copy-local-multidevice.c)

Wait for Completion

- Since kernel launching is non-blocking. we need to wait for them to complete.
- We simply use clWaitForEvents, and wait for all events to complete.

Wait for Kernels

Example 13: (matrixMul-time-copy-local-multidevice.c)

```
clWaitForEvents(DEVICENUM, events);
printf("Kernel execution completes.\n");
```

Discussion

How many events do we need to wait for the whole thing to complete?

Kernel

- The kernel function is very similar to the previous local memory version.
- We also use the constant DEVICENUM to denote the number of devices.

Kernel

Example 14: (mul-local-multidevice-kernel.cl)

```
1 #define N 1024
2 #define Blk 64
3 #define DEVICENUM 2
4 #define BSIDE (N / Blk)
```

Kernel

Example 15: (mul-local-multidevice-kernel.cl)

```
6
   __kernel void mul(__global int A[N/DEVICENUM][N],
                      __global int B[N][N],
8
                      __global int C[N/DEVICENUM][N])
9
   {
10
     int globalRow = get_global_id(0);
11
     int globalCol = get_global_id(1);
12
     int localRow = get_local_id(0);
13
     int localCol = get_local_id(1);
14
15
     __local int ALocal[BSIDE][BSIDE];
16
     __local int BLocal[BSIDE][BSIDE];
```

NDRange

- We declare A and C as N / DEVICENUM by N arrays, because the domain is partitioned among devices.
- The rest of the code is not changed at all!

Reason

- The reason for this "no need to change" is that a kernel function is acting from a local view, i.e., it is a work item within a work group, so all the operations are still the same.
- The kernel still thinks from the point of view of blocks, and the related operations are done in local indices.
- ullet We only need the global index while accessing A, B, and C.
- Since the kernel is given the corresponding parts of A and B, so the operation is correct.

Discussion

• Convince yourself that the code is correct.



Timing

- We would like to know the detailed timing of the kernels.
- In particular we would like to prove the two kernel runs in parallel.
- To do so we need the absolute time of the events from all devices.

Base

- The time returned by the event is difficult to interpret.
- We would like to establish a relative time for inspection.
- We choose the time the first kernel entering the queue as the base time, and report the relative time to it for ease understanding.

Example 16: (matrixMul-time-copy-local-multidevice.c)

- We loop through all devices and get timing information from all of them.
- This part is the same as before.

154

155

156

157

158

159

160

161

162

163

164

165

166

Example 17: (matrixMul-time-copy-local-multidevice.c)

```
for (int device = 0; device < DEVICENUM; device++) {</pre>
  cl_ulong timeEnterQueue, timeSubmit, timeStart,
    timeEnd;
  status =
    clGetEventProfilingInfo(events[device],
      CL_PROFILING_COMMAND_QUEUED,
      sizeof(cl_ulong), &timeEnterQueue, NULL);
  assert(status == CL_SUCCESS);
  status =
    clGetEventProfilingInfo(events[device],
      CL_PROFILING_COMMAND_SUBMIT,
      sizeof(cl_ulong), &timeSubmit, NULL);
  assert(status == CL_SUCCESS);
```

168

170

171

176

177

Example 18: (matrixMul-time-copy-local-multidevice.c)

```
status =
169
           clGetEventProfilingInfo(events[device],
             CL_PROFILING_COMMAND_START,
             sizeof(cl_ulong), &timeStart, NULL);
172
        assert(status == CL_SUCCESS);
173
        status =
174
           clGetEventProfilingInfo(events[device],
175
             CL_PROFILING_COMMAND_END,
              sizeof(cl_ulong), &timeEnd, NULL);
        assert(status == CL_SUCCESS);
```

Relative Time

- In addition to queuing time, submission time, and execution, we also report the relative time when the four events happened.
- The relative time is in nanosecond.

Print Time

Example 19: (matrixMul-time-copy-local-multidevice.c)

```
179
        printf("\nkernel entered queue at %f\n",
180
                (timeEnterQueue - base) / NANO2SECOND);
181
        printf("kernel submitted to device at %f\n",
182
                (timeSubmit - base) / NANO2SECOND);
183
        printf("kernel started at %f\n",
184
                (timeStart - base) / NANO2SECOND):
185
        printf("kernel ended at %f\n",
186
                (timeEnd - base) / NANO2SECOND);
187
        printf("kernel queued time %f seconds\n",
188
                (timeSubmit - timeEnterQueue) / NANO2SECOND);
189
        printf("kernel submission time %f seconds\n",
190
                (timeStart - timeSubmit) / NANO2SECOND);
191
        printf("kernel execution time %f seconds\n",
192
              (timeEnd - timeStart) / NANO2SECOND);
      }
193
```

Demonstration

 Run the matrixMul-time-copy-local-multidevice-cl program.



Discussion

• Does the kernels run in parallel? Observe the results and give your answer.

Dependency

- We may launch multiple kernels into multiple devices.
- Kernels may have dependency, and we need to ensure dependency among kernels.

Example

- In the following example we compute G = (A + B) + (D + E), where all vector has length N.
- We compute C = A + B, then F = D + E, then we compute G = C + F. The last computation must wait for the first two to complete.
- We will use one device for each of the three additions.

Constants

• We first declare the constants and variables.

Example 20: (vectorAdd-dependency.c)

```
#define CL_USE_DEPRECATED_OPENCL_2_O_APIS
   #include <stdio.h>
3
   #include <assert.h>
4
   #include <CL/cl.h>
5
   #define N (65536 * 4)
6
   #define MAXGPU 10
   #define MAXK 1024
8
   #define MAXI.OG 4096
9
   #define DEVICENUM 3
10
   #define NANO2SECOND 100000000.0
11
12
   cl_uint A[N], B[N], C[N], D[N], E[N], F[N], G[N];
```

Initialization

• We initialize the A, B, D and E vectors.

Example 21: (vectorAdd-dependency.c)

Buffers

- We then create buffers for kernel parameters.
- A and B are read only, and need to be copied into device.
- C is both read and write enable.

Example 22: (vectorAdd-dependency.c)

```
cl mem bufferA =
 clCreateBuffer(context,
                 CL MEM READ ONLY | CL MEM COPY HOST PTR.
                 N * sizeof(cl_uint), A, &status);
assert(status == CL_SUCCESS);
cl mem bufferB =
 clCreateBuffer(context.
                 CL MEM READ ONLY | CL MEM COPY HOST PTR.
                 N * sizeof(cl uint). B. &status):
assert(status == CL_SUCCESS);
cl mem bufferC =
 clCreateBuffer(context.
                 CL_MEM_READ_WRITE,
                 N * sizeof(cl uint), NULL, &status):
assert(status == CL_SUCCESS);
```

88

89

90

91

92

93

94

95

96

97

98

99

100

101

102

Buffers

• Similarly we create buffers for D, E and F.

Example 23: (vectorAdd-dependency.c)

```
cl mem bufferD =
 clCreateBuffer(context,
                 CL MEM READ ONLY | CL MEM COPY HOST PTR.
                 N * sizeof(cl_uint), D, &status);
assert(status == CL_SUCCESS);
cl mem bufferE =
 clCreateBuffer(context.
                 CL MEM READ ONLY | CL MEM COPY HOST PTR.
                 N * sizeof(cl uint), E, &status):
assert(status == CL_SUCCESS);
cl mem bufferF =
 clCreateBuffer(context,
                 CL_MEM_READ_WRITE,
                 N * sizeof(cl uint), NULL, &status):
assert(status == CL_SUCCESS);
```

104

105

106

107

108

109

110

111

112

113

114

115 116

117

118

Buffers

- We create buffers for G.
- This buffer is write only and use the host memory directly, so the host can print it directly.

Example 24: (vectorAdd-dependency.c)

120

121

122

123 124

125

Discussion

• What are the characteristic for read only, read and write, and write only buffers respectively?

Buffers

- Both NDRange and the work group are one dimensional.
- The size of NDRange is N.
- The size of a work group is 256.

Example 25: (vectorAdd-dependency.c)

```
size_t globalThreads[] = {(size_t)N};
size_t localThreads[] = {256};
```

$$C = A + B$$

- Now we launch the first kernel for C = A + B.
- We have already built three command queues for three devices.
- We use commandQueue[0] for the first addition.
- We also declared three events to denote the completion of the three kernels that we will launch.

Example 26: (vectorAdd-dependency.c)

```
130
       cl event events[3]:
131
       status = clSetKernelArg(kernel, 0, sizeof(cl_mem),
132
                                (void*)&bufferA):
133
       assert(status == CL_SUCCESS);
134
       status = clSetKernelArg(kernel, 1, sizeof(cl_mem),
135
                                (void*)&bufferB):
136
       assert(status == CL_SUCCESS);
137
       status = clSetKernelArg(kernel, 2, sizeof(cl_mem),
138
                                (void*)&bufferC):
139
       assert(status == CL_SUCCESS);
140
141
       status =
142
         clEnqueueNDRangeKernel(commandQueue[0], kernel, 1, NULL,
143
                                 globalThreads, localThreads,
144
                                 0, NULL, &(events[0]));
145
       assert(status == CL SUCCESS):
```

$$F = D + E$$

- Similarly we launch the second kernel for F = D + E.
- Note that we used commandQueue[1] and events[1].

Example 27: (vectorAdd-dependency.c)

```
147
       status = clSetKernelArg(kernel, 0, sizeof(cl_mem),
148
                                (void*)&bufferD):
149
       assert(status == CL SUCCESS):
150
       status = clSetKernelArg(kernel, 1, sizeof(cl_mem),
151
                                (void*)&bufferE):
152
       assert(status == CL SUCCESS):
153
       status = clSetKernelArg(kernel, 2, sizeof(cl_mem),
154
                                (void*)&bufferF):
155
       assert(status == CL SUCCESS):
156
       status =
157
         clEnqueueNDRangeKernel(commandQueue[1], kernel, 1, NULL,
158
                                 globalThreads, localThreads,
159
                                 0, NULL, &(events[1]));
160
       assert(status == CL SUCCESS):
```

$$G = C + F$$

- Finally we launch the third kernel for G = C + F.
- Note that we used commandQueue[2] and events[2].

Example 28: (vectorAdd-dependency.c)

```
162
       status = clSetKernelArg(kernel, 0, sizeof(cl_mem),
163
                                (void*)&bufferC);
164
       assert(status == CL SUCCESS):
165
       status = clSetKernelArg(kernel, 1, sizeof(cl_mem),
166
                                (void*)&bufferF):
167
       assert(status == CL SUCCESS):
168
       status = clSetKernelArg(kernel, 2, sizeof(cl_mem),
169
                                (void*)&bufferG):
170
       assert(status == CL SUCCESS):
171
       status =
172
         clEnqueueNDRangeKernel(commandQueue[2], kernel, 1, NULL,
173
                                 globalThreads, localThreads,
174
                                 0, NULL, &(events[2]));
175
       assert(status == CL SUCCESS):
```

Wait for Events

• We wait for all three kernel to finish.

Example 29: (vectorAdd-dependency.c)

```
status = clWaitForEvents(DEVICENUM, events);
assert(status == CL_SUCCESS);
printf("All three kernels complete.\n");
```

177

178

179

Wait for Events

- The rest of the code just prints the timing information, checks for correctness, and releases resources.
- Please refer to previous discussion.

Demonstration

• Run the vectorAdd-dependency-cl.

Discussion

- Does the program produce the correct answer?
- If not what is the possible reason?

Problem

- The previous program waits for all three kernels to complete, but the third kernel did *not* wait for the first two.
- We will use the wait for events mechanism to launch the third kernel.

Prototype 30: clEnqueueNDRangeKernel.h

```
cl int
   clEnqueueNDRangeKernel (cl_command_queue command_queue,
3
                             cl_kernel kernel,
4
                             cl uint work dim.
5
                             const size_t *global_work_offset,
6
                             const size_t *global_work_size,
7
                             const size t *local work size.
8
                             cl_uint num_events_in_wait_list,
9
                             const cl_event *event_wait_list,
10
                             cl_event *event);
```

$$G = C + F$$

- We launch the third kernel and wait for the first two events in the events array.
- Set num_events_in_wait_list to 2 and event_wait_list to events.

Example 31: (vectorAdd-dependency-correct.c)

```
162
       status = clSetKernelArg(kernel, 0, sizeof(cl_mem),
163
                                (void*)&bufferC):
164
       assert(status == CL SUCCESS):
165
       status = clSetKernelArg(kernel, 1, sizeof(cl_mem),
166
                                (void*)&bufferF):
167
       assert(status == CL SUCCESS):
168
       status = clSetKernelArg(kernel, 2, sizeof(cl_mem),
169
                                (void*)&bufferG):
170
       assert(status == CL SUCCESS):
171
       status =
172
         clEnqueueNDRangeKernel(commandQueue[2], kernel, 1, NULL,
173
                                 globalThreads, localThreads,
174
                                 2, events, &(events[2]));
175
       assert(status == CL SUCCESS):
```

Demonstration

• Run the vectorAdd-dependency-correct-cl.

Discussion

- Does the program produce the correct answer?
- Observe the timing and make sure the first two kernels run in parallel, and the third kernel will wait for the first two.

flFinish

- The host can also explicitly wait for the first two kernels to finish before launching the third kernel by clFinish.
- We just need to wait for the command queues to the first two devices to become empty.

Example 32: (vectorAdd-dependency-clfinish.c)

```
clFinish(commandQueue[0]);
clFinish(commandQueue[1]);
status = clSetKernelArg(kernel, 0, sizeof(cl_mem),
                        (void*)&bufferC):
assert(status == CL_SUCCESS);
status = clSetKernelArg(kernel, 1, sizeof(cl_mem),
                        (void*)&bufferF):
assert(status == CL SUCCESS):
status = clSetKernelArg(kernel, 2, sizeof(cl_mem),
                        (void*)&bufferG):
assert(status == CL SUCCESS):
status =
 clEnqueueNDRangeKernel(commandQueue[2], kernel, 1, NULL,
                         globalThreads, localThreads,
                         0, NULL, &(events[2]));
assert(status == CL SUCCESS):
```

162

163

164

165

166

167

168

169

170

171

172

173

174

175

176

177

Demonstration

• Run the vectorAdd-dependency-clfinish-cl program.

Discussion

- Does the program produce the correct answer?
- Observe the timing and make sure the first two kernels run in parallel, and the third kernel will wait for the first two.
- Compare the timing results between vectorAdd-dependency-correct-cl and vectorAdd-dependency-clfinish-cl, especially in when the last kernel joined the command queue.

Group Size

- The previous matrix multiplication program has extremely good performance due to two reasons.
 - It uses local memory to speed up data access.
 - It has a larger group size (256).
- Now we would like to study the effects of group size on performance.

Compute Units

- A device has only a limited number of compute units.
- The work items of a work group will occupy a compute unit.
- For a given number of work items, if the work group size is small, then the number of work groups becomes large, and some work groups will wait for compute unit.
- In other words, a small work group size limits the the number of work items that we can process in parallel.

Classroom

- A school has only a limited number of classrooms.
- The students are divided into group, and one group will occupy a classroom.
- For a given number of students, if the group size is small, then the number of groups becomes large, and some groups will wait for classroom.
- In other words, a small group size limits the the number of students that can study in parallel.

Discussion

• Give your own example of this phenomenon.

Compute Unit Number

- We need only one GPU device.
- We call clGetDeviceInfo with CL_DEVICE_MAX_COMPUTE_UNITS to get the number of compute units on GPU.
- We need the number of compute units to determine the work group size.

Example 33: (vectorAdd-groupsize.c)

```
23
     cl_device_id GPU[MAXGPU];
24
     cl_uint GPU_id_got;
25
     status =
26
        clGetDeviceIDs(platform_id, CL_DEVICE_TYPE_GPU,
27
                       MAXGPU, GPU, &GPU_id_got);
28
     assert(status == CL_SUCCESS && GPU_id_got >= 1);
29
     printf("There are %d GPU devices\n", GPU_id_got);
30
     cl_uint unit;
31
     status =
32
       clGetDeviceInfo(GPU[0], CL_DEVICE_MAX_COMPUTE_UNITS,
33
                        sizeof(cl_uint), &unit, NULL);
34
     assert(status == CL_SUCCESS);
35
     printf("# of compute units is %d\n", unit);
```

Group Size

- We fix the number of work items to *N*, and vary the size of the work group from 1 to 256.
- We wait for the event before launching the next kernel.
- We will record the timing information into a file vectorAdd-groupsize.dat.

Example 34: (vectorAdd-groupsize.c)

```
108
       size_t workItem[] = {(size_t)N};
109
       FILE *timefp = fopen("vectorAdd-groupsize.dat", "w");
110
       assert(timefp != NULL);
111
       for (int groupSize = 1; groupSize <= 256; groupSize *= 2)</pre>
112
         cl event event:
113
         size t localSize[1]:
114
         localSize[0] = groupSize;
115
         status =
116
           clEnqueueNDRangeKernel (commandQueue, kernel, 1, NULL,
117
                                   workItem, localSize,
118
                                   O, NULL, &event);
119
         assert(status == CL_SUCCESS);
120
         /* waitforevent */
121
         status = clWaitForEvents(1, &event);
122
         assert(status == CL_SUCCESS);
123
         printf("The kernel with group size %d completes.\n",
124
                groupSize);
```

Get Timing Information

- The rest of the code will retrieve timing information from the event.
- We output the group size and the execution time into a file vectorAdd-groupsize.dat.

Example 35: (vectorAdd-groupsize.c)

```
printf("\nkernel entered queue at %f\n",
         (timeEnterQueue - base) / NANO2SECOND);
  printf("kernel submitted to device at %f\n".
         (timeSubmit - base) / NANO2SECOND);
  printf("kernel started at %f\n",
         (timeStart - base) / NANO2SECOND):
  printf("kernel ended at %f\n",
         (timeEnd - base) / NANO2SECOND):
  printf("kernel queued time %f seconds\n",
         (timeSubmit - timeEnterQueue) / NANO2SECOND);
  printf("kernel submission time %f seconds\n".
         (timeStart - timeSubmit) / NANO2SECOND);
  printf("kernel execution time %f seconds\n",
         (timeEnd - timeStart) / NANO2SECOND):
  fprintf(timefp, "%d %f\n", groupSize,
          (timeEnd - timeStart) / NANO2SECOND);
}
```

157

158

159

160

161

162

163

164

165

166

167

168

169

170

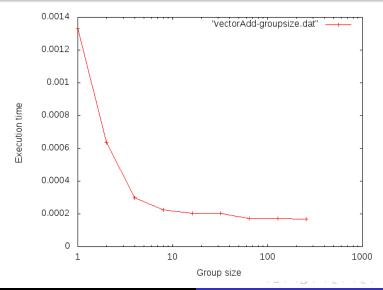
171

172 173

Demonstration

- We plot the kernel execution time for different group sizes.
- The x-xis is in log scale.

Execution Time



Discussion

• Observe the execution time and draw you conclusion about the effects of group size on performance.