

# Advanced OpenMP Programming

Pangfeng Liu  
National Taiwan University

May 9, 2015

- We will observe the timing when threads access variables at different locations.
- The intuition is that local variables are much easier to access than the global ones, and we need to confirm that.

# Compute Square Roots

- Write a program that computes the square root from 0 to 999999999 and assign the value to a variable  $v$ .
- We first observe the timing when we place the variable  $v$  at global area.

## Example 1: (assign.c)

```
1  #include <omp.h>
2  #include <stdio.h>
3  #include <math.h>
4  #define N 999999999
5  int main()
6  {
7      double v;
8      #pragma omp parallel for
9          for (int i = 1; i <= N; i++)
10         v = sqrt(i);
11     return 0;
12 }
```

# Demonstration

- Run the assign-uni program.
- Run the assign-omp program.

# Discussion

- Compare the execution time of `assign-uni` and `assign-omp`.

# Private Variable

- We now declare  $v$  as private and observe the timing.

## Example 2: (assign-private.c)

```
1  #include <omp.h>
2  #include <stdio.h>
3  #include <math.h>
4  #define N 999999999
5  int main()
6  {
7      double v;
8      #pragma omp parallel for private(v)
9          for (int i = 1; i <= N; i++)
10         v = sqrt(i);
11     return 0;
12 }
```



# Demonstration

- Run the assign-private-omp program.

# Discussion

- Compare the execution time of `assign-private-omp` with previous `assign-uni` and `assign-omp`.
- What is the reason for this performance difference?

# Heap Variable

- We now put  $v$  into the heap and observe the timing.

### Example 3: (assign-heap.c)

```
1  #include <omp.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <math.h>
5  #define N 999999999
6  int main()
7  {
8      double *v = malloc(sizeof(double));
9      #pragma omp parallel for
10     for (int i = 1; i <= N; i++)
11         *v = sqrt(i);
12     return 0;
13 }
```

# Demonstration

- Run the assign-heap-omp program.

# Discussion

- Compare the execution time of `assign-heap-omp` with previous programs.
- What is the reason for this performance difference?

# Prime Number Counting

- We want to count the number of prime numbers.
- We start with an array of numbers, and assuming that every number is a prime number.
- We start with the smallest prime number in the array, and mark *every* multiple of it as *composite* (non-prime number).
- We repeat this process until no new prime numbers are found.

### Example 4: (prime.c)

```
1  #include <stdio.h>
2  #include <math.h>
3  #include <stdlib.h>
4  #include <assert.h>
5  #include "omp.h"
6
7  #define N 400000000
8
9  char notPrime[N];
10 int nPrime = 0;
11 char primes[N];
```



```
13 int main(int argc, char *argv[])
14 {
15     assert(argc == 2);
16     int n = atoi(argv[1]);
17     int bound = round(sqrt(n));
18
19     for (int i = 2; i <= bound; i++)
20         if (!notPrime[i]) {
21 #pragma omp parallel for
22         for (int j = 2 * i; j < n; j += i)
23             notPrime[j] = 1;
24     }
```

```
26     int nPrime = 0;
27 #pragma omp parallel for reduction(+ : nPrime)
28     for (int i = 2; i < n; i++)
29         if (notPrime[i] == 0)
30             nPrime++;
31
32     printf("number of prime is %d\n", nPrime);
33     return 0;
34 }
```

# Variables

- Array `notPrime` keeps track of the status of a number. If we know a number `i` is *not* prime, we set `notPrime[i]` to 1.
- The prime numbers will be kept in another array `primes`.
- The range to be tested (`n`) is given as a command line argument.

# Try All Possibilities

- Try all numbers from 2 to  $\sqrt{n}$ .
- If  $i$  is a prime number, mark the all multiple of  $i$  as *not* prime.
- It is obvious that the first for loop *cannot* be parallelized because of dependency, so we parallelize the second for loop.

# Reduction

- The number of prime number can be obtained by counting the number of zeros in array `notPrime`.
- We use a reduction on the variable `nPrime` to simplify the process.

# Demonstration

- Run the prime-uni program.
- Run the prime-omp program.

# Discussion

- Compute the speedup of the previous prime counting program.
- Is there any optimization that can improve the performance?

# Efficiency

- The previous program parallelized the inner for loop only.
- The previous program will go through multi-threading every time a prime number is found.
- This incurs overheads of creating and destroying the threads.



# Spawn and Join



- Single thread for the outer loop.
- Multi-thread for the inner loop.

# parallel Once

- We would like to avoid the overheads in creating and destroying threads, so we put a `parallel` in front of the first for loop.
- The *entire* two level loop is run by all threads.
- Then we share the workload of the second for loop to improve performance, since most work is done in the second loop.

**Example 5: (prime-inner.c)**

```
12 int main(int argc, char *argv[])
13 {
14     assert(argc == 2);
15     int n = atoi(argv[1]);
16     int bound = round(sqrt(n));
17
18     #pragma omp parallel
19         for (int i = 2; i <= bound; i++)
20             if (!notPrime[i])
21                 #pragma omp for
22                     for (int j = 2 * i; j < n; j += i)
23                         notPrime[j] = 1;
```

# Efficient

- All threads run the first loop, and share the second loop.
- It is OK for all threads to run the first for loop simultaneously, since they will synchronize at every second loop.
- Index variables *i* and *j* are private.

# Spawn and Join



- Multi-thread for both outer and inner loops.
- The work in the outer loop is duplicated, but it is more efficient than creating and joining threads.

# Demonstration

- Run the prime-uni program.
- Run the prime-omp program.
- Run the prime-inner-omp program.

# Discussion

- Compare the execution time of all three prime counting programs.
- Is there any further optimization that can improve the performance?

# Synchronization

- A `for` pragma will synchronize all threads before leaving the `for` statement.
- If the `for` statement is still within the same `parallel` directive, then a barrier synchronization will do.
- If the `for` statement is at the end of the `parallel` directive, threads will be joined by the master thread and destroyed.



# nowait

- Sometimes we do not wish the the threads to wait for each other.
- Those finish earlier can go on to the next statement to improve performance.
- We can only do this if the following statement does *not* depend on the previous statement.
- In this case we can use `nowait` clause.

# nowait

```
1  nowait
```

## Two loops

- We place two for directives with a parallel directive.
- The first for has an ascending workload, and the second loop has a descending workload.

**Example 6: (2for.c)**

```
6  int main(int argc, char *argv[])
7  {
8      assert(argc == 3);
9      omp_set_num_threads(atoi(argv[1]));
10     int n = atoi(argv[2]);
11     printf("# of proc = %d\n", omp_get_num_procs());
12     printf("# of loop iterations = %d\n", n);
```

```
14     double t = omp_get_wtime();
15     #pragma omp parallel
16     {
17     #pragma omp for
18         for (int i = 0; i < n; i++)
19             sleep(i);
20     #pragma omp for
21         for (int i = n - 1; i >= 0; i--)
22             sleep(i);
23     }
24     printf("time = %f\n", omp_get_wtime() - t);
```

# Demonstration

- Run the `2for-omp` program with 4 threads and 8 iterations, and observe the timing.

# Discussion

- Describe the the execution time of `2for-omp` and make sure that it is reasonable.

# Wait

- In the previous program If the work of the second loop does not depend on the first loop, then we can let the threads go to the second directly.
- Since the first for has an *ascending* workload, and the second loop has a *descending* workload, if we let the threads to go to the second loop then the workload imbalance will be reduced.
- We only need to add a `nowait` clause at the first for directive.



**Example 7: (2for-nowait.c)**

```
14  double t = omp_get_wtime();
15  #pragma omp parallel
16  {
17  #pragma omp for nowait
18      for (int i = 0; i < n; i++)
19          sleep(i);
20  #pragma omp for
21      for (int i = n - 1; i >= 0; i--)
22          sleep(i);
23  }
24  printf("time = %f\n", omp_get_wtime() - t);
```

# Demonstration

- Run the `2for-nowait-omp` program with 4 threads and 8 iterations, and observe the timing.

# Discussion

- Describe the the execution time of `2for-omp` and make sure that it is reasonable.

# nowait

- We consider our previous prime number counting program.
- Previously the second marking `for` will synchronize before going back to the outside `for`.
- We would like to remove this synchronization, and let each threads to start the “prime” finding as soon as possible.
- This will not cause a race condition since the threads will still synchronize at the beginning of the next inner loop.

**Example 8: (prime-inner-nowait.c) nowait**

```
18 #pragma omp parallel
19     for (int i = 2; i <= bound; i++)
20         if (!notPrime[i])
21 #pragma omp for nowait
22     for (int j = 2 * i; j < n; j += i)
23         notPrime[j] = 1;
```

# Demonstration

- Run `prime-inner-omp` and `prime-inner-nowait-omp` and compare their execution time.

# Discussion

- Why can the `nowait` clause improve performance?

# Compute $\pi$

- We calculate  $\pi$  by integrating  $f(x) = \frac{4}{1+x^2}$ , where  $x$  is from 0 to 1.
- Divide the interval into  $N$  pieces, and assume the area in each interval is a trapezoid, then sum the area of these  $N$  trapezoids into a variable area.
- Note that in each interval you only need to compute  $\frac{4}{1+x^2}$  once.



# Parallel Version

- We use  $x$  to denote the  $x$  coordinate, and  $area$  for the area in the integral.
- We parallelize the program by adding `parallel` for `pragma` and declare  $x$  as `private`.

Example 9: (pi.c) Compute  $\pi$ 

```
1  #include <omp.h>
2  #include <stdio.h>
3  #define N 1000000000
4  int main()
5  {
6      double x;
7      double area = 0.0;
8      double t = omp_get_wtime();
9      #pragma omp parallel for private(x)
10     for (int i = 0; i < N; i++) {
11         x = (i + 0.5) / N;
12         area += 4.0/(1.0 + x * x);
13     }
14     double pi = area / N;
15     t = omp_get_wtime() - t;
16     printf("execution time is %f\n", t);
17     printf("pi = %f\n", pi);
18     return 0;
19 }
```

# Demonstration

- Run `pi-omp`.

# Discussion

- Is the answer from the parallel version correct?
- Find out why the sequential version will not compile and fix it.  
Then compare the the execution time of these two programs.

# Atomic

- We did not declare `area` private because it has the final global answer.
- As a result the operation on `area` must be *atomic* to avoid race condition.
- We simply use `critical` directive on the loop body to ensure atomic condition.

Example 10: (pi-critical.c) Compute  $\pi$ 

```
8  double x;  
9  double area = 0.0;  
10 double t = omp_get_wtime();  
11 #pragma omp parallel for private(x)  
12   for (int i = 0; i < N; i++)  
13 #pragma omp critical  
14   {  
15       x = (i + 0.5) / N;  
16       area += 4.0 / (1.0 + x * x);  
17   }  
18 double pi = area / N;  
19 t = omp_get_wtime() - t;  
20 printf("execution time is %f\n", t);  
21 printf("pi = %f\n", pi);
```

# Demonstration

- Run the `pi-critical-omp` program.

# Discussion

- Is the answer from the parallel version correct?
- Compare the the execution time of with the two previous programs.



# Atomic

- It appears that we do not need to make the entire loop body critical because `x` is already private.
- We now only use critical directive on the area summation.

# Critical Section

- It is essential to reduce the size of the critical section,
- We want a thread to get through a critical section as soon as possible, so that other threads can get into the critical section as well.

**Example 11: (pi-critical-small.c) Compute  $\pi$** 

```
8  double x;  
9  double area = 0.0;  
10 double t = omp_get_wtime();  
11 #pragma omp parallel for private(x)  
12   for (int i = 0; i < N; i++)  
13     {  
14       x = (i + 0.5) / N;  
15 #pragma omp critical  
16       area += 4.0 / (1.0 + x * x);  
17     }  
18 double pi = area / N;  
19 t = omp_get_wtime() - t;  
20 printf("execution time is %f\n", t);  
21 printf("pi = %f\n", pi);
```

# Demonstration

- Run the `pi-critical-small-omp` program.

# Discussion

- Is the answer from the parallel version correct?
- Compare the the execution time of with the three previous programs.

# Private Variables

- The performance improvement is very limited by a smaller critical section because it is not very different from the previous implementation.
  - Two statements v.s. one statement.
- If there are many statements in the loop body the benefit will more obvious.

# Critical Section

- The number of critical sections is enormous.
- We would like to remove these time consuming critical sections by letting all threads to work on its integrals by summing into its own area.
- The idea is to use a global array to store the individual area, then the master threads can add sum them up.

# Array Implementation

- We need an array for threads to store its area.
- During each iteration a thread needs to call `omp_get_thread_num()` to know where to store its area, which is a significant overheads.



**Example 12: (pi-array.c)**

```
7  double x;  
8  double area[MAXT] = {0.0};  
9  double t = omp_get_wtime();  
10 #pragma omp parallel for private(x)  
11   for (int i = 0; i < N; i++) {  
12       x = (i + 0.5) / N;  
13       area[omp_get_thread_num()] +=  
14         4.0 / (1.0 + x * x);  
15   }  
16   t = omp_get_wtime() - t;  
17   double areaSum = 0.0;  
18   for (int i = 0; i < omp_get_num_procs(); i++)  
19       areaSum += area[i];  
20   double pi = areaSum / N;
```

# Demonstration

- Run the `pi-array-omp` program.

# Discussion

- Is the answer from the parallel version correct?
- Compare the the execution time of with the previous programs.
- Is there any way to reduce the overheads in calling `omp_get_thread_num()`?

# Reduction

- We now use a reduction to compute the integral.
- The implementation is much cleaner and (hopefully) with better performance because it has been optimized by the OpenMP library.

**Example 13: (pi-reduction.c) Compute  $\pi$** 

```
6  double x;  
7  double area = 0.0;  
8  double t = omp_get_wtime();  
9  #pragma omp parallel for private(x) \  
10     reduction(+ : area)  
11     for (int i = 0; i < N; i++) {  
12         x = (i + 0.5) / N;  
13         area += 4.0 / (1.0 + x * x);  
14     }  
15     double pi = area / N;  
16     t = omp_get_wtime() - t;  
17     printf("execution time is %f\n", t);  
18     printf("pi = %f\n", pi);
```

# Demonstration

- Run the `pi-reduction-omp` program.

# Discussion

- Is the answer from the parallel version correct?
- Compare the the execution time of with the previous programs.

# Game of Life

- A two-dimensional board with cells. A cell could either *live* or *dead*.
- The status of a cell evolves according to its status and the status of its eight neighbors.
- A dead cell with exactly *three* live neighbors becomes a live cell.
- A live cell with two or three live neighbors stays live, otherwise it become dead.



# Double Buffer

- It is intuitive that we keep the status of cells in a two dimensional array.
- Since the status of cell depends on the status of others, if we modify a cell directly it will affect the computation on other cells. This causes “race” condition during the update.
- Instead we use *two* arrays  $A$  and  $B$  to store the status of cells. In initially the status is in  $A$ .

# Iterations

- We repeat the following steps.
  - We set the cell status of  $B$  according to  $A$  in the first, third, etc. iterations.
  - We set the cell status of  $A$  according to  $B$  in the second, fourth, etc. iterations.

# Double Buffers

- We use two buffers A and B to avoid data inconsistency in updating the status of cells.
- We use a macro to compute the number of live neighboring cells – a live cell has 1 and a dead cell has 0.
- We pad the broad boundary so that we can use a single macro to compute the number of live neighboring cells

**Example 14: (life.c)**

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <omp.h>
4
5  #define MAXN 4096
6  #define SIDE (MAXN + 2)
7
8  #define nLiveNeighbor(A, i, j) \
9      A[i + 1][j] + A[i - 1][j] + A[i][j + 1] + \
10     A[i][j - 1] + A[i + 1][j + 1] + A[i + 1][j - 1] + \
11     A[i - 1][j + 1] + A[i - 1][j - 1]
12
13 char A[SIDE][SIDE];
14 char B[SIDE][SIDE];
```

# Print

- We use a simple printing routine to print the status of cells.

```
16 void print(char A[SIDE][SIDE], int n)
17 {
18     for (int i = 1; i <= n; i++) {
19         for (int j = 1; j <= n; j++)
20             printf("%2d ", A[i][j]);
21         printf("\n");
22     }
23 }
```

# Input

- if the flag `READINPUT` is set the main program will read the size of the board, the number of generations, and the cell status from `stdin`.
- Otherwise it will generate a random input of size 4096 by 4096, and repeat for ten generations.

```
25 int main()
26 {
27     int n, generation, cell;
28     #ifdef READINPUT
29         scanf("%d%d", &n, &generation);
30         for (int i = 1; i <= n; i++)
31             for (int j = 1; j <= n; j++) {
32                 scanf("%d", &cell);
33                 A[i][j] = cell;
34             }
35     #else
36         n = 4096;
37         generation = 20;
38         for (int i = 1; i <= n; i++)
39             for (int j = 1; j <= n; j++)
40                 A[i][j] = rand() % 2;
41     #endif
```



# Dead or Alive

- Depending on the generation the program will set the cell status of B with A, or in the other direction.
- A cell will be alive only if it is dead now and has three live neighbors, or it is live and it has two or three live neighbors now.
- We simply use a `parallel for` directive to distribute the workload.

```
43     int nln;  
44     for (int g = 0; g < generation; g++)  
45         if (g % 2 == 0)  
46             #pragma omp parallel for          /* from A to B */  
47                 for (int i = 1; i <= n; i++)  
48                     for (int j = 1; j <= n; j++) {  
49                         nln = nLiveNeighbor(A, i, j);  
50                         B[i][j] = ((A[i][j] == 0 && nln == 3) ||  
51                             (A[i][j] == 1 && (nln == 2 || nln == 3)));  
52                     }  
53         else  
54             #pragma omp parallel for          /* from B to A */  
55                 for (int i = 1; i <= n; i++)  
56                     for (int j = 1; j <= n; j++) {  
57                         nln = nLiveNeighbor(B, i, j);  
58                         A[i][j] = ((B[i][j] == 0 && nln == 3) ||  
59                             (B[i][j] == 1 && (nln == 2 || nln == 3)));  
60                     }
```

# Final Status

- if the flag `PRINT` is set the main program will output the final cell status.

```
62 #ifdef PRINT
63     if (generation % 2 == 0)
64         print(lifeA, n);
65     else
66         print(lifeB, n);
67 #endif
68     return 0;
69 }
```

# Demonstration

- Run the `life-uni` program.
- Run the `life-omp` program.

# Discussion

- Compute the speedup of the parallel program.
- What kind of scheduling policy was used?
- Is there any optimization that can improve the performance?